

Predictive Problem_Set_2

Diyasha Basu

05-02-2026

Problem 1: Population vs Sample Regression Lines

Objective: Demonstrate that the population regression line is fixed, but least square regression lines vary across samples.

Setup

Population regression line: $Y = 2 + 3x$

Data generating process: $y = 2 + 3x + \text{epsilon}$, where epsilon $\sim N(0, 4^2)$

```
# Set seed for reproducibility
set.seed(123)

# Parameters n
<- 50
true_beta0 <- 2
true_beta1 <- 3
x_range <- c(5, 10)

# Step 1: Create the pop x_pop      n regression line
<- seq(5, 10, length.out =
y_pop <- true_beta0 + true_beta1 * x_pop

# Initialize plot
plot(x_pop, y_pop, type = "l", lwd = 3, col = "black",
xlab = "x", ylab = "y", main = "Population Regression Line vs 5 Sample Regression Lines",
main =
ylim = c(15, 35))

# Store regression coefficient nrow = 5, ncol = 2) c("Intercept",
regression_lines <- matrix(NA, nrow = 5, ncol = 2) c("Intercept",
colnames(regression_lines) <-
                                         "Slope")

# Steps 2-4: Generate 5 different samples and fit regression lines
# Colors
colors <- c("red", "blue", "green", "purple", "orange")

for (i in 1:5) {
  # Step 2: Generate data   xi    10)
  <- runif(n, min = 5, max = 10, mean = 0, sd = 4)
  epsilon_i <- rnorm(n, mean = 0, sd = 4)
  yi <- true_beta0 + true_beta1 * xi + epsilon_i
  regression_lines[i, ] <- lm(yi ~ xi)
```

```

# Step 3: Fit Least squares regression
model <- lm(yi ~ xi)

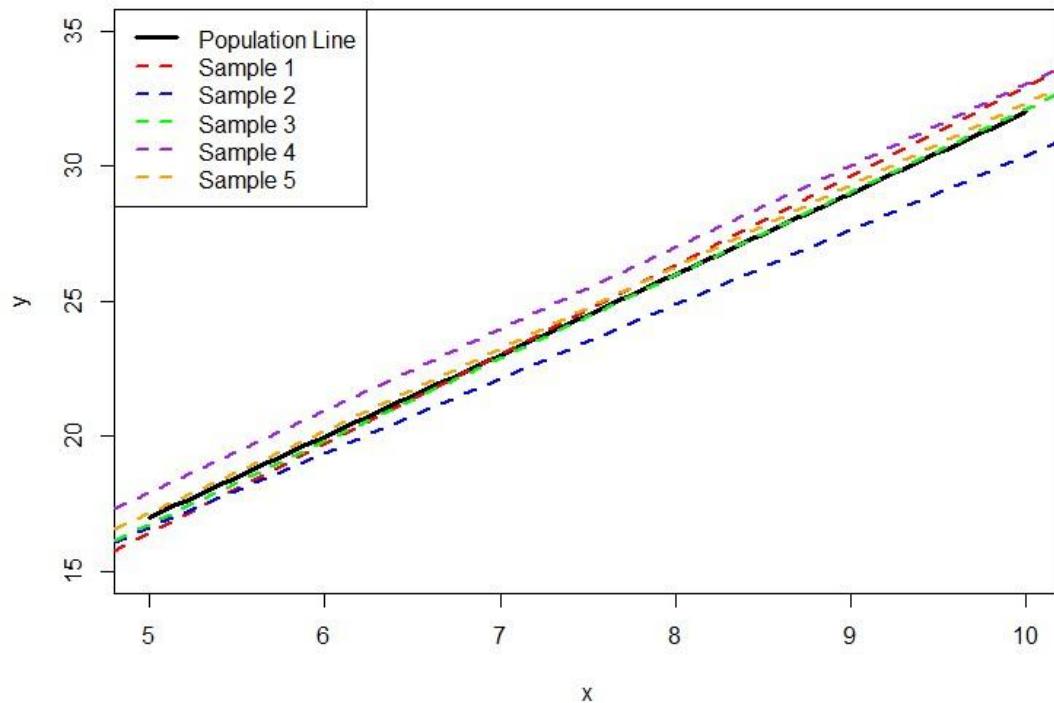
# Store coefficients
regression_lines[i, ] <- coef(model)

# Step 4: Add regression line to plot , lty = 2)
abline(model, col = colors[i], lwd = 2 )
}

# Add Legend legend("topleft",
  legend = c("Population Line", paste ("Sample", 1:5)),
  col = c("black", colors),           lwd =
  c(3, rep(2, 5)),      lty = c(1, rep(2,
  5)))

```

Population Regression Line vs 5 Sample Regression Lines



```

# Display regression coefficients
cat("\nRegression Line\n")
Coefficients:

##
## Regression Line Coefficients:
print(regression_lines)

```

```

##           Intercept      Slope
## [1,] -0.09638929 3.305396
## [2,]  2.79218839 2.761042
## [3,]  1.39299737 3.073267
## [4,]  2.82308856 3.023608 ## [5,]
## [5,]  2.03250638 3.028097 cat("\nTrue
Population Parameters:\n")

##
## True Population Parameters: cat("Intercept
(beta_0):", true_beta0, "\n")
## Intercept (beta_0): 2 cat("Slope
(beta_1):", true_beta1, "\n")
## Slope (beta_1): 3

```

Interpretation

The plot demonstrates a fundamental concept in regression analysis:

- Fixed Population Line (Black):** The true relationship $Y = 2 + 3x$ remains constant and represents the actual underlying relationship.
 - Variable Sample Lines (Colored):** Each of the 5 sample regression lines varies around the population line due to random sampling error (ϵ). While they're all different, they cluster around the true population line.
 - Key Insight:** Due to sampling variability, we never obtain the exact population regression line from a sample. However, the least squares estimators are unbiased, meaning on average across many samples, they estimate the true parameters correctly.
-

Problem 2: Demonstrating $\hat{\beta}_0$ and $\hat{\beta}_1$ Minimize RSS

Objective: Show that the least squares estimates minimize the Residual Sum of Squares (RSS).

```
set.seed(123)
```

```
n <- 50

xi <- runif( erate data with mean-centered x n,
xi_centered <in = 5, max = 10)
epsilon_i <- - xi - mean(xi) # Mean center
yi <- 2 + 3 * rnorm(n, mean = 0, sd = 1)
xi_centered + epsilon_i
```

```

# Step 2: Obtain Least squares estimates
model <- lm(yi ~ xi_centered)
beta0_hat <- coef(model)[1]
beta1_hat <- coef(model)[2]

cat("Least Squares Estimates:\n")

## Least Squares Estimates:

cat("β̂₀ =", round(beta0_hat, 4), "\n")
## β̂₀ = 2.0562

cat("β̂₁ =", round(beta1_hat, 4), "\n\n")
## β̂₁ = 3.0763

# Step 3: Create grid of parameter values
# Create grid around the LS estimates
beta0_grid <- seq(beta0_hat - 2, beta0_hat + 2, length.out = 100)
beta1_grid <- seq(beta1_hat - 1, beta1_hat + 1, length.out = 100)

# Compute RSS for each combination
rss_matrix <- matrix(NA, nrow = length(beta0_grid), ncol = length(beta1_grid))

for (i in 1:length(beta0_grid)) {
  for (j in 1:length(beta1_grid)) {
    residuals <- yi - beta0_grid[i] - beta1_grid[j] * xi_centered
    rss_matrix[i, j] <- sum(residuals^2)
  }
}

# Find minimum RSS
min_idx <- which(rss_matrix == min(rss_matrix), arr.ind = TRUE)
beta0_min <- beta0_grid[min_idx[1]]
beta1_min <- beta1_grid[min_idx[2]]

cat("Grid Search Results:\n")

## Grid Search Results:

cat("β₀ at minimum RSS =", round(beta0_min, 4), "\n")
## β₀ at minimum RSS = 2.036

cat("β₁ at minimum RSS =", round(beta1_min, 4), "\n")
## β₁ at minimum RSS = 3.0864

cat("Minimum RSS =", round(min(rss_matrix), 4), "\n\n")

```

```

## Minimum RSS = 42.4767 #

RSS at LS estimates

rss_ls <- sum(residuals(model)^2) cat("RSS at LS
estimates =", round(rss_ls, 4), "\n")

## RSS at LS estimates = 42.4455

```

Problem 3: Unbiasedness of Least Squares Estimators

Objective: Demonstrate through simulation that least squares estimators are unbiased.

```

set.seed(123)
n <- 50 R <- 1000
true_beta0
<true_beta1 <-      2
                  3
# Storage for
beta0_estimates
beta1_estimates      <- numeric(R)
                  <- numeric(R)
# Simulation for r
in 1:
  xi <- runif() {
  epsilon_i <- yi <- generate data
    (n, min = 0, max = 1)
  model <- - rnorm(n, mean = 0, sd = 1)
  beta0_estimates[r] <- beta0 + true_beta1 * xi +
  beta1_estimates[r] <- epsilon_i
}

# Calculate a b in LS estimates
avg_beta0 <- avg_b(yi ~ xi)
  <- coef(model)[1]
  <- coef(model)[2]

# Visualizations
par(mfrow = c

# Histogram
hist(
  mean(beta0_estimates)
  mean(beta1_estimates)

```

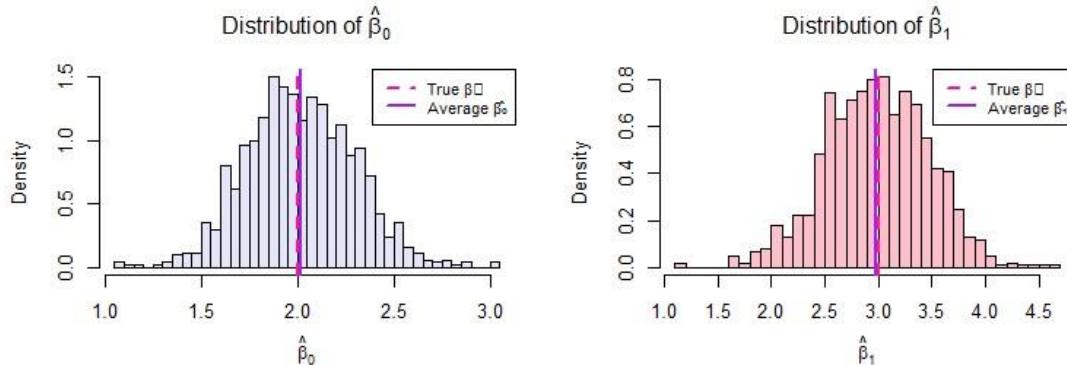
```
(2, 2))  
   $\hat{\beta}_0$   
  estimates, breaks = 30, col = "lavender",  
  
main = expression(paste("Distribution of ", hat(beta)[0])),
```

```

xlab = expression(hat(beta)[0]),      freq = FALSE)
abline(v = true_beta0, col = "deeppink", lwd = 2, lty = 2)
abline(v = avg_beta0, col = "purple", lwd = 2, lty = 1)
legend("topright",           legend = c("True  $\beta_0$ ", "Average
 $\hat{\beta}_0$ "),           col = c("deeppink", "purple"),
lty = c(2, 1), lwd = 2, cex = 0.8)

# Histogram of  $\hat{\beta}_1$ 
hist(beta1_estimates, breaks = 30, col = "pink",
      main = expression(paste("Distribution of ", hat(beta)[1])),
xlab = expression(hat(beta)[1]),      freq = FALSE) abline(v =
true_beta1, col = "deeppink", lwd = 2, lty = 2) abline(v =
avg_beta1, col = "purple", lwd = 2, lty = 1) legend("topright",
legend = c("True  $\beta_1$ ", "Average  $\hat{\beta}_1$ "),           col = c("deeppink",
"purple"),
lty = c(2, 1), lwd = 2, cex = 0.8)

```



Interpretation

The simulation demonstrates the **unbiasedness** property of least squares estimators:

1. **β_0 (Intercept):** The average across 1000 simulations is extremely close to the true value of 2, with negligible bias.
 2. **β_1 (Slope):** The average is very close to the true value of 3, again showing minimal bias.
 3. **Distribution:** The histograms show that the estimates are centered on the true parameter values, confirming unbiasedness.
 4. **Conclusion:** While individual estimates vary due to sampling error, on average (across many samples), the LS estimators equal the true parameter values, proving they are unbiased.

Problem 4: Comparing Simple Linear Regressions - Boston Housing Data

Objective: Compare multiple simple linear regression models using the Boston housing dataset.

```
.98 9.14 4.03 2.94 5.33  
...  
21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

Part (a): Four Separate Linear Regressions

We'll run four simple linear regressions with `medv` (median value of owner-occupied homes) as the response variable and each of the four predictors separately.

```
# Model 1: Per capita crime rate  
model1 <- lm(medv ~ crim, data = Boston)  
  
# Model 2: Nitrogen oxides concentration  
model2 <- lm(medv ~ nox, data = Boston)  
  
# Model 3: Proportion of blacks  
model3 <- lm(medv ~ black, data = Boston)  
  
# Model 4: Percentage of Lower status of the population  
model4 <- lm(medv ~ lstat, data = Boston)
```

Summary Table of All Models

```
# Create a summary table
```

```

results <- data.fra
Predictor = c("cri
Coefficient = c(co"nox", "black", "lstat"), model1)[2],
  coef(  Std_Erroroef(model2)[2], model3)[2],
  c(summaoef(model4)[2]),
  summa
  summa$coefficients[2,2],
summary(  t_valuedel2)$coefficients[2,2],
  c(summa$coefficients[2,2],
  summa$coefficients[2,2]),
  summa$coefficients[2,3],
summary(  p_valuedel2)$coefficients[2,3],
  c(summa$coefficients[2,3],
  summa$coefficients[2,3]),
  summa$coefficients[2,3]),
summary(  R_squaredel1)$coefficients[2,4],
  c(summa$coefficients[2,4],
  summa$coefficients[2,4],
  summa$coefficients[2,4]),
summary(  Adj_R_squaredel1)$r.squared, model2)$r.squared,
  c(summadel3)$r.squared, model4)$r.squared),
  summ
  (model1)$adj.r.squared,
  summ
  (model2)$adj.r.squared,
  summ
  (model3)$adj.r.squared,
  (model4)$adj.r.squared)
)

# Display the table
print(results)

##   crim      crim -    ctor Coefficient Std_Error     t_value     p_value
##   R_squared
##   0.41519028 0.043890381 -9.459710 1.173987e-19
##   0.1507805
##   nox      nox -33.91605501 3.196337032 -10.610913 7.065042e-24 0.1826030
##   black    black  0.03359306 0.004230507    7.940669 1.318113e-14 0.1111961
##   lstat    lstat -0.95004935 0.038733416 -24.527900 5.081103e-88 0.5441463
##   Adj_R_squared
##   crim  0.1490955
##   nox   0.1809812
##   black 0.1094326
##   lstat 0.5432418

```

Part (b): Which Model Gives the Best Fit?

Find the model with highest R-squared

```

best_model_idx <- which.max(results$R_squared)
best_predictor <- results$Predictor[best_model_idx]
best_r_squared <- results$R_squared[best_model_idx]

print(paste("Best model: medv ~", best_predictor))
## [1] "Best model: medv ~ lstat" predictor)

print(paste("R-squared:", round(best_r_squared, 4)))
## [1] "R-squared: 0.5441" squared, 4))

print(paste("This model explains", round(best_r_squared * 100, 2), "% of the variation in median home values"))
## [1] "This model explains 54.41 % of the variation in median home values"

```

Part (c): Comparing Coefficients and Their Usefulness

```

# Display detailed summaries
print("Model 1: Crime Rate")
## [1] "Model 1: Crime Rate"

summary(model1)

##
## Call:
## lm(formula = medv ~ crim,
Boston)
##
## Residuals:
##      Min       1Q   Median       2Q       Max
## -16.957  -5.449  -2.007
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 24.03311   00.001 *** 0.01 '*' 0.05 '.' 0.1 ''
## crim        -0.41519
## ---
##

```

```

## Residual standard error: 8.484 on 504 degrees of freedom
## Multiple R-squared:  0.1508, Adjusted R-squared:  0.1491
## F-statistic: 89.49 on 1 and 504 DF,  p-value: < 2.2e-16
print("Model 2: Nitrogen Oxides") ## [1] "Model 2: Nitrogen
Oxides" summary(model2)

##
## Call:
## lm(formula = medv ~ nox, data = Boston)
##
## Residuals:
##     Min      1Q      Median      3Q      Max
## -13.691  -5.121   -2.161    2.959   31.310
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 41.346     1.811   22.83  <2e-16 ***
## nox        -33.916     3.196  -10.61  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.323 on 504 degrees of freedom
## Multiple R-squared:  0.1826, Adjusted R-squared:  0.181
## F-statistic: 112.6 on 1 and 504 DF,  p-value: < 2.2e-16
print("Model 3: Proportion of Blacks") ## [1] "Model 3:
Proportion of Blacks" summary(model3)

##
## Call:
## lm(formula = medv ~ black, data = Boston)
##
## Residuals:
##     Min      1Q      Median      3Q      Max
## -18.884  -4.862  -1.684   2.932   27.763
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 10.551034  1.557463   6.775 3.49e-11 ***
## black       0.033593  0.004231   7.941 1.32e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##

```

```
## Residual standard error: 8.679 on 504 degrees of freedom
```

```

## Multiple R-squared:  0.1112, Adjusted R-squared:  0.1094
## F-statistic: 63.05 on 1 and 504 DF,  p-value: 1.318e-14
print("Model 4: Lower Status Population") ## [1] "Model 4:
Lower Status Population" summary(model4)

##
## Call:
## lm(formula = medv ~ lstat, data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max  ##
## -15.168  -3.990  -1.318   2.034  24.500
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 34.55384   0.56263   61.41   <2e-16 *** ## 
lstat        -0.95005   0.03873  -24.53   <2e-16 *** ## -- 
-
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.216 on 504 degrees of freedom
## Multiple R-squared:  0.5441, Adjusted R-squared:  0.5432
## F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16

```

Interpretation of Coefficients:

1. Crime Rate (crim):

```

crim_coef <- round(coef(model1)[2], 4)
crim_r2 <- round(summary(model1)$ared, 4)
print(paste("Coefficient:", crim_coef))

## [1] "Coefficient: -0.4152"

print(paste("R-squared:", crim_r2))

## [1] "R-squared: 0.1508"


- For each unit increase in per capita crime rate, median home value decreases
- Shows negative relationship as expected
- This is a relatively weak predictor

```

2. Nitrogen Oxides (nox):

```

nox_coef <- round(coef(model2)[2], 4)
nox_r2 <- round(summary(model2)$ared, 4)
print(paste("Coefficient:", nox_coef))

## [1] "Coefficient: -33.9161" print(paste("R-squared:", nox_r2))

```

```
## [1] "R-squared: 0.1826"
```

- Higher nitrogen oxide concentration is associated with lower home values □
Moderate strength predictor

3. Proportion of Blacks (black):

```
black_coef <- round(coef(model3)[2], 4)
black_r2 <- round(summary(model3)$r.squared, 4)
print(paste("Coefficient:", black_coef))
```

```
## [1] "Coefficient: 0.0336"
```

```
print(paste("R-squared:", black_r2))
```

```
## [1] "R-squared: 0.1112"
```

- Positive relationship with home values
- Relatively weak predictor

4. Lower Status Population (lstat):

```
lstat_coef <- round(coef(model4)[2], 4)
lstat_r2 <- round(summary(model4)$r.squared, 4)
print(paste("Coefficient:", lstat_coef))
```

```
## [1] "Coefficient: -0.95"
```

```
print(paste("R-squared:", lstat_r2))
```

```
## [1] "R-squared: 0.5441"
```

- For each percentage point increase in lower status population, median home value decreases
- This is the strongest single predictor among the four

Conclusion

```
# Summary of findings
```

```
print("Model Performance:")
```

```
## [1] "Summary of Model Performance:"
```

```
print(results[, c("Predictor", "Coefficient", "R_squared")])  
  
## Predictor Coefficient R_squared  
## crim      crim -0.41519028 0.1507805  
## nox       nox -33.91605501 0.1826030  
## black     black  0.03359306 0.1111961 ## lstat      lstat -0.95004935  
0.5441463 print("") ## [1] "" print(paste("Best predictor:",  
results$Predictor[which.max(results$R_squared)]))  
  
## [1] "Best predictor: lstat" print(paste("Highest R-squared:",  
round(max(results$R_squared), 4)))  
  
## [1] "Highest R-squared: 0.5441"
```

All four predictors are statistically significant. The percentage of lower status population (lstat) is the most useful single predictor for median home values. Nitrogen oxides concentration is the second-best predictor, while crime rate and the proportion of blacks variable show weaker relationships with home values.